

---

Ravi Shripad

Roll no.58

## ✓ House Price prediction using Linear Regression - SingleVarile

```
1 import pandas as pd
2 from sklearn.linear_model import LinearRegression
3 import matplotlib.pyplot as plt
```

```
➞ C:\ProgramData\Anaconda3\lib\site-packages\scipy\__init__.py:146: UserWarning: A NumPy version
  warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")
```

## ✓ Load Dataset

```
1 Dataset = pd.read_csv('house dataset.csv')
2 Dataset.head()
```

```
➞
```

	area	price
0	8450	208500
1	9600	181500
2	11250	223500
3	9550	140000
4	14260	250000

## ✓ Load Summariz

```
1 print(Dataset.shape)
2 print(Dataset.head(5))
```

```
➞ (1460, 2)
```

	area	price
0	8450	208500
1	9600	181500
2	11250	223500
3	9550	140000
4	14260	250000

```
1 Dataset.info()
```

```

↳ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  ---
0   area     1460 non-null    int64
1   price    1460 non-null    int64
dtypes: int64(2)
memory usage: 22.9 KB

```

```
1 Dataset.describe()
```

```

↳

```

	area	price
count	1460.000000	1460.000000
mean	10516.828082	180921.195890
std	9981.264932	79442.502883
min	1300.000000	34900.000000
25%	7553.500000	129975.000000
50%	9478.500000	163000.000000
75%	11601.500000	214000.000000
max	215245.000000	755000.000000

## Visualize Dataset

```


1 plt.xlabel('Area')
2 plt.ylabel('price')
3 plt.scatter(Dataset.area, Dataset.price, color='blue', marker='*')

```

```

↳ <matplotlib.collections.PathCollection at 0x14af2ff11f0>

```



## Segreate Dataset into Input X & Output Y

```
1 X = Dataset.drop('price', axis='columns')
2 X
```

```
↔
```

	area
0	8450
1	9600
2	11250
3	9550
4	14260
...	...
1455	7917
1456	13175
1457	9042
1458	9717
1459	9937

1460 rows × 1 columns

```
1 Y = Dataset.price
2 Y
```

```
↔
```

0	208500
1	181500
2	223500
3	140000
4	250000
...	...
1455	175000
1456	210000
1457	266500
1458	142125
1459	147500

Name: price, Length: 1460, dtype: int64

## ✓ Training Dataset using Linear Regression

```
1 model = LinearRegression()
2 model.fit(X,Y)
```

```
↔ LinearRegression()
```

## ✓ Predicted Price for Land sq.Feet of custom values

```

1 x=int(input('Enter house Squar fit'))
2 LandAreainSqFt=[[x]]
3 PredictedmodelResult = model.predict(LandAreainSqFt)
4 print(PredictedmodelResult)

```

```

Enter house Squar fit3000
[165136.067752]
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have v:
warnings.warn(

```

## Checking model is right

## Theory Calculation

✓  $Y=m*X+b$  (m is coefficient and b is intercept)

Coefficient -m

```

1 m=model.coef_
2 print(m)

```

```
[2.09997195]
```

intercept - b

```

1 b=model.intercept_
2 print(b)

```

```
158836.1518968766
```

✓  $Y=mx+b$

x is independent variable- input - area

```

1 y = m*x + b
2 print("The price of {0} Squar feet Land is: {1}".format(x,y[0]))

```

```
The price of 3000 Squar feet Land is: 165136.06775199962
```

✓ Part B-Exam marks

```
1 import pandas as pd
2 from sklearn.linear_model import LinearRegression
```

```
1 dataset = pd.read_csv('exam data.csv')
```

```
1 dataset.head(10)
```



	hours	age	internet	marks
0	6.83	15	1	78.50
1	6.56	16	0	76.74
2	NaN	17	1	78.68
3	5.67	18	0	71.82
4	8.67	19	1	84.19
5	7.55	20	0	81.18
6	6.67	15	0	76.99
7	8.99	16	0	85.46
8	5.19	17	1	70.66
9	6.75	18	0	77.82

```
1 print(dataset.shape)
2 print(dataset.head(5))
```



```
(201, 4)
   hours  age  internet  marks
0   6.83   15         1   78.50
1   6.56   16         0   76.74
2    NaN   17         1   78.68
3   5.67   18         0   71.82
4   8.67   19         1   84.19
```

```
1 X = dataset.iloc[:, :-1].values
2 print(X.shape)
3 X
```



```

[ 7.7, 19. , 1. ],
[ 7.3 , 20. , 0. ],
[ 7.79, 15. , 1. ],
[ 6.75, 16. , 0. ],
[ 7.87, 17. , 1. ],
[ 5.38, 18. , 0. ],
[ 7.8 , 19. , 1. ],
[ 5.07, 20. , 0. ],
[ 7.95, 15. , 1. ],
[ 8.35, 16. , 0. ],
[ 5.19, 17. , 0. ],
[ 7.19, 18. , 0. ],
[ 7.35, 19. , 1. ],
[ 5.22, 20. , 1. ],
[ 5.39, 15. , 1. ],
[ 5.39, 16. , 1. ],
[ 8.93, 17. , 1. ],
[ 5.79, 18. , 0. ],
[ 8.42, 19. , 1. ],
[ 7.26, 20. , 0. ],
[ 6.97, 15. , 1. ],
[ 5.55, 16. , 1. ],
[ 8.66, 17. , 0. ],
[ 8.61, 18. , 1. ],
[ 5.22, 19. , 1. ],
[ 8.05, 20. , 0. ],
[ 8.87, 15. , 1. ],
[ 5.54, 16. , 0. ],
[ nan, 17. , 0. ],
[ 7.26, 18. , 1. ],
[ 5.79, 19. , 0. ],
[ 5.22, 20. , 0. ],
[ 8.71, 15. , 1. ],
[ 7.55, 16. , 1. ],
[ 6.35, 17. , 1. ],
[ 7.53, 18. , 0. ],
[ 8.56, 19. , 1. ],
[ 8.94, 20. , 1. ],
[ 6.6 , 15. , 1. ],
[ 8.35, 16. , 1. ],
[ 4.15, 15. , 0. ]])

```

```
1 dataset.columns[dataset.isna().any()]
```

```
➦ Index(['hours'], dtype='object')
```

```
1 dataset.hours = dataset.hours.fillna(dataset.hours.mean())
```

```
1 X = dataset.iloc[:, :-1].values
```

```
2 print(X.shape)
```

```
3 X
```

```
➦
```

```
[ 0.95      , 19.      , 1.      ],
[ 5.14      , 20.      , 0.      ],
[ 8.39      , 15.      , 1.      ],
[ 6.18      , 16.      , 0.      ],
[ 7.53      , 17.      , 1.      ],
[ 7.86      , 18.      , 0.      ],
[ 7.7       , 19.      , 1.      ],
[ 7.3       , 20.      , 0.      ],
[ 7.79      , 15.      , 1.      ],
[ 6.75      , 16.      , 0.      ],
[ 7.87      , 17.      , 1.      ],
[ 5.38      , 18.      , 0.      ],
[ 7.8       , 19.      , 1.      ],
[ 5.07      , 20.      , 0.      ],
[ 7.95      , 15.      , 1.      ],
[ 8.35      , 16.      , 0.      ],
[ 5.19      , 17.      , 0.      ],
[ 7.19      , 18.      , 0.      ],
[ 7.35      , 19.      , 1.      ],
[ 5.22      , 20.      , 1.      ],
[ 5.39      , 15.      , 1.      ],
[ 5.39      , 16.      , 1.      ],
[ 8.93      , 17.      , 1.      ],
[ 5.79      , 18.      , 0.      ],
[ 8.42      , 19.      , 1.      ],
[ 7.26      , 20.      , 0.      ],
[ 6.97      , 15.      , 1.      ],
[ 5.55      , 16.      , 1.      ],
[ 8.66      , 17.      , 0.      ],
[ 8.61      , 18.      , 1.      ],
[ 5.22      , 19.      , 1.      ],
[ 8.05      , 20.      , 0.      ],
[ 8.87      , 15.      , 1.      ],
[ 5.54      , 16.      , 0.      ],
[ 6.98142857, 17.      , 0.      ],
[ 7.26      , 18.      , 1.      ],
[ 5.79      , 19.      , 0.      ],
[ 5.22      , 20.      , 0.      ],
[ 8.71      , 15.      , 1.      ],
[ 7.55      , 16.      , 1.      ],
[ 6.35      , 17.      , 1.      ],
[ 7.53      , 18.      , 0.      ],
[ 8.56      , 19.      , 1.      ],
[ 8.94      , 20.      , 1.      ],
[ 6.6       , 15.      , 1.      ],
[ 8.35      , 16.      , 1.      ],
[ 4.15      , 15.      , 0.      ]])
```

1 dataset.hours

```
0      6.830000
1      6.560000
2      6.981429
3      5.670000
4      8.670000
...
196    8.560000
197    8.940000
198    6.600000
199    8.350000
200    4.150000
Name: hours, Length: 201, dtype: float64
```

```
1 Y = dataset.iloc[:, -1].values
2 Y
```

```
➦ array([78.5 , 76.74, 78.68, 71.82, 84.19, 81.18, 76.99, 85.46, 70.66,
        77.82, 75.37, 83.88, 79.5 , 80.76, 83.08, 76.03, 76.04, 85.11,
        82.5 , 80.58, 82.18, 83.36, 70.67, 75.02, 70.96, 83.33, 74.75,
        75.65, 74.15, 80.17, 82.27, 76.14, 71.1 , 84.35, 83.08, 76.76,
        81.24, 78.21, 73.08, 83.23, 70.27, 86.41, 71.1 , 82.84, 82.38,
        72.96, 77.46, 70.11, 72.38, 71.41, 72.22, 77.77, 84.44, 71.45,
        82.21, 85.48, 75.03, 86.65, 70.9 , 71.7 , 73.61, 79.41, 76.19,
        80.43, 85.78, 70.06, 81.25, 81.7 , 69.27, 82.79, 71.8 , 71.79,
        74.97, 78.61, 77.59, 72.33, 72.08, 77.33, 70.05, 73.34, 84. ,
        82.93, 76.63, 75.36, 77.29, 72.87, 73.4 , 81.74, 71.85, 84.6 ,
        79.56, 82.1 , 72.08, 79.1 , 81.01, 76.48, 75.39, 68.57, 83.64,
        82.3 , 75.18, 82.03, 82.99, 79.26, 77.55, 77.07, 72.1 , 73.25,
        74.25, 70.58, 81.08, 75.04, 76.38, 80.86, 78.42, 74.44, 70.34,
        85.04, 73.61, 75.55, 76.2 , 82.69, 76.83, 79.53, 83.57, 85.95,
        76.02, 77.65, 77.01, 74.49, 73.19, 71.86, 75.8 , 72.46, 78.39,
        83.48, 83.15, 71.22, 85.98, 83.91, 84.58, 80.31, 82.55, 75.52,
        83.82, 85.15, 82.75, 74.34, 82.02, 86.12, 71.87, 76.7 , 81.7 ,
        70.78, 78.45, 70.2 , 83.37, 75.52, 81.57, 80.72, 80.81, 79.49,
        79.17, 77.07, 82.04, 71.94, 81.6 , 70.79, 82.68, 83.08, 71.18,
        77.63, 77.78, 70.4 , 73.02, 71.11, 85.96, 73.64, 84.24, 78.17,
        77.19, 71.83, 86.99, 83.87, 71.5 , 79.63, 85.1 , 72.01, 77.27,
        79.87, 73.14, 70.51, 84.03, 79.64, 74.24, 81.67, 84.68, 86.75,
        78.05, 83.5 , 81.45])
```

```
1 model=LinearRegression()
2 model.fit(X,Y)
```

```
➦ LinearRegression()
```

```
1 a=[[1,75,1]]
2 predicedModelresult = model.predict(a)
3 print(PredictedmodelResult)
```

```
➦ [165136.067752]
```

1 Start coding or [generate](#) with AI.

1 Start coding or [generate](#) with AI.